

Generative AI Tools Comparison: Diffusion Models vs GANs and Tool Analysis

Raphael Adesta Pratidina - GAIS Assignment 1 Part 3

October 2025 Semester

Course Overview

This document provides a comprehensive comparison of generative AI tools, focusing on:

- **Core architectures:** Diffusion models vs GANs
- **Tool comparison:** Learning Assistant (GPT-5) vs Qwen Chat
- **Technical breakdowns:** Generation processes, architectures, and limitations

1 Generative AI Fundamentals

1.1 Diffusion Models vs GANs: Core Concepts

1.1.1 Diffusion Models

Key Takeaway

Definition: Probabilistic models that generate images through iterative denoising of random noise.

Analogy: Like solving a jigsaw puzzle where you start with a completely scrambled image and gradually reveal the correct pieces.

Technical Definition:

$$\text{Diffusion} = \text{Forward Process} \circ \text{Reverse Process}$$

Where:

- Forward process: Gradually adds Gaussian noise to real data
- Reverse process: Learns to remove noise to reconstruct original data

Crucial Insight

Why Diffusion?

- **Stable training:** No adversarial instability issues
- **High fidelity:** Better at capturing fine details
- **Controllable generation:** Strong text-to-image alignment

Disadvantage: Slower generation due to iterative process

1.1.2 Generative Adversarial Networks (GANs)

Key Takeaway

Definition: Two neural networks (generator/discriminator) trained adversarially.

Analogy: Like a game of cat-and-mouse where the generator tries to fool the discriminator.

Technical Definition:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Crucial Insight

Why GANs?

- **Fast generation:** Single forward pass
- **Photorealism:** Excellent for style transfer
- **Simpler architecture:** Easier to implement

Disadvantage: Training instability (mode collapse, vanishing gradients)

1.2 Real-World Applications

Remember

Diffusion Model Example: Stable Diffusion

- **Use case:** Creative design, marketing, game art
- **Strengths:**
 - High detail preservation
 - Strong text alignment
 - Versatile prompt handling

Remember

GAN Example: StyleGAN

- **Use case:** Synthetic data generation, avatar creation
- **Strengths:**
 - Photorealistic outputs
 - Fast generation
 - Style transfer capabilities

Feature	Learning Assistant (GPT-5)	Qwen Chat
Primary Model	GPT-5 (mini/nano)	Qwen-3 series (Max, VL, Coder)
Modalities	Text-only	Multimodal (text, image, video, audio)
Use Case	Module-specific Q&A	General-purpose AI assistant
Access	School-restricted	Publicly available

2 Tool Comparison: Learning Assistant vs Qwen Chat

2.1 Tool Overview

2.2 Architectural Breakdown

2.2.1 Learning Assistant Architecture

Code Structure

Transformer Decoder Architecture

- **Input:** Tokenized text prompt
- **Embedding Layer:** Converts tokens to vectors
- **Transformer Layers:**
 - Self-attention mechanism
 - Position-wise feedforward networks
- **Output:** Probability distribution over vocabulary

2.2.2 Qwen Chat Architecture

Code Structure

Multimodal Architecture

- **Text Path:**
 - Tokenization → Embedding → Transformer
- **Image Path:**
 - Vision Transformer (ViT) for feature extraction
 - Cross-modal fusion via attention
- **Specialized Models:**
 - Qwen-Image: Diffusion-based image generation
 - Qwen-Image-Edit: Latent diffusion for editing
 - Qwen-3 Omni: Multimodal fusion

2.3 Performance Comparison

Remember

Accuracy Comparison

- Both models provide accurate technical explanations
- Qwen offers more detailed explanations with better prompt flexibility
- Learning Assistant shows more concise, structured responses

Key Takeaway

Response Quality Metrics

- **Learning Assistant:**
 - Consistent structure
 - Limited creativity
 - Higher hallucination rate without tools
- **Qwen Chat:**
 - Flexible response styles
 - Lower hallucination with RAG/web search
 - Higher creativity potential

3 Technical Deep Dive

3.1 Diffusion Model Generation Process




Figure 1: Diffusion Model Generation Pipeline

3.1.1 Line-by-Line Diffusion Process

Code Documentation

Text Encoding Function

```
def encode_prompt(prompt: str) -> torch.Tensor:  
    """  
    Encodes text prompt into semantic embedding  
  
    Parameters:  
    prompt (str): Input text description  
    Returns:  
    torch.Tensor: Semantic embedding vector (d_model dimensions)  
    """  
    # 1. Tokenize input text  
    tokens = tokenizer(prompt, return_tensors="pt")  
  
    # 2. Pass through text encoder (CLIP)  
    embeddings = text_encoder(tokens)  
  
    # 3. Project to latent space dimension  
    latent_embedding = projection_layer(embeddings)  
  
    return latent_embedding
```

Key Parameters:

- `d_model`: Latent space dimension (typically 768-1024)
- `tokenizer`: Specialized tokenizer for text-to-vector conversion
- `projection_layer`: Linear layer to match latent space dimensions

Why This Design?

- **Advantage:** Captures semantic meaning beyond surface text
- **Disadvantage:** Computationally expensive for long prompts

3.2 GAN Training Process




Figure 2: GAN Adversarial Training Loop

3.2.1 Key Components Analysis

Crucial Insight

Why U-Net for Diffusion?

- **Architectural choice:** U-Net's skip connections preserve spatial information
- **Advantage over GANs:** Better at capturing fine details through hierarchical feature learning
- **Implementation detail:**
 - Downsampling path: Extracts context
 - Upsampling path: Reconstructs details
 - Skip connections: Preserve high-resolution features

Crucial Insight

Why Cross-Attention in Diffusion?

- **Purpose:** Aligns generated image with text prompt
- **Implementation:**
 - Text embeddings condition each denoising step
 - Attention weights determine which image regions to modify
- **Advantage:** Enables precise control over generated content
- **Disadvantage:** Computationally expensive ($O(n^2)$ complexity)

4 Implementation Details

4.1 Common Architectural Components

4.1.1 Batch Normalization

Remember

Definition: Normalizes layer inputs to have zero mean and unit variance

Purpose in Diffusion Models:

- Stabilizes training by reducing internal covariate shift
- Helps gradient flow through deep U-Net architectures

Why Used Here?

- Diffusion models often use very deep U-Nets (12+ layers)
- BatchNorm helps prevent vanishing gradients in deep networks

Implementation Details:

- Applied after each convolutional layer
- Uses moving averages for statistics during inference

4.1.2 LeakyReLU Activation

Remember

Definition: Variant of ReLU that allows small negative values

Mathematical Form:

$$\text{LeakyReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

where $\alpha \approx 0.01$

Purpose in GANs:

- Prevents "dead neurons" problem in discriminator
- Helps with gradient flow for negative inputs

Why Used Here?

- Standard ReLU can cause discriminator to ignore negative samples
- LeakyReLU maintains gradient flow for all inputs

4.1.3 BCE Loss (Binary Cross Entropy)

Remember

Definition: Measures difference between binary predictions and true labels

Mathematical Form:

$$\text{BCE}(y, \hat{y}) = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Purpose in GANs:

- Discriminator uses BCE to distinguish real/fake
- Generator uses BCE to maximize discriminator's error

Why Used Here?

- More stable than MSE for binary classification
- Better handles class imbalance (real vs fake)

4.2 Optimization Techniques

Crucial Insight

Why Adam Optimizer?

- **Advantage:**
 - Adaptive learning rates per parameter
 - Combines momentum and RMSprop
 - Works well with batch sizes ≥ 32
- **Implementation:**
 - Default parameters: $\beta_1 = 0.9$, $\beta_2 = 0.999$
 - Learning rate typically 10^{-4} to 10^{-3}
- **Why Not SGD?**
 - GAN training benefits from adaptive learning rates
 - SGD often requires careful tuning of momentum

5 Tool-Specific Analysis

5.1 Learning Assistant Analysis

5.1.1 Strengths

Remember

- **Domain specialization:** Pre-loaded with module-specific knowledge
- **Structured responses:** Consistent format for educational content
- **Guardrails:** Strong ethical compliance for school environment

5.1.2 Limitations

Warning

- **No multimodality:** Text-only input/output
- **Limited context:** Fixed context window (typically 4096 tokens)
- **No web search:** Relies solely on pre-loaded knowledge
- **Hallucination risk:** Without retrieval tools, prone to factual errors

5.2 Qwen Chat Analysis

5.2.1 Strengths

Remember

- **Multimodal capabilities:** Handles text, image, video, audio
- **Tool integration:** Web search, RAG, code execution
- **Prompt flexibility:** Adaptable response styles
- **Open models:** Ability to run locally for privacy

5.2.2 Limitations

Warning

- **Censorship:** Strict guardrails on sensitive topics
- **Performance variability:** Model quality depends on variant selection
- **Data collection:** Free version collects user data
- **Latency:** Slower response times with multimodal inputs




Stage	Dimensions
Text Embedding	768
Latent Space	$64 \times 64 \times 4$
Final Image	$512 \times 512 \times 3$

Figure 3: Diffusion Model Generation Pipeline with Dimensions

6 Workflow Diagrams

6.1 Diffusion Model Generation Flow

6.2 GAN Training Loop



Batch Size: 64

Figure 4: GAN Adversarial Training Loop with Weight Updates

6.3 End-to-End System Flow




Figure 5: End-to-End GAN Training System Flow

7 Common Implementation Errors

8 Extensions and Advanced Topics

8.1 Advanced Diffusion Techniques

8.1.1 Latent Diffusion Models

Crucial Insight

Why Latent Space?

- **Advantage:**
 - Reduces computational cost by working in compressed space
 - Preserves high-quality generation
 - Enables faster sampling (50-100 steps vs thousands)
- **Implementation:**
 - Use autoencoder to compress images to latent space
 - Apply diffusion in latent space
 - Decode final latent representation
- **Example:** Stable Diffusion uses VAE (Variational Autoencoder)

8.1.2 ControlNet

Crucial Insight

Enhanced Control

- **Purpose:** Adds conditional control to diffusion models
- **Implementation:**
 - Additional encoder network processes input conditions
 - Condition information fused via attention
 - Supports multiple control types:
 - * Edge maps
 - * Depth maps
 - * Pose estimation
 - * Keypoints
- **Example Use Cases:**
 - Sketch-to-image generation
 - Pose-preserving image editing
 - Style transfer with structural constraints

8.2 Advanced GAN Techniques

8.2.1 StyleGAN Architecture

Crucial Insight

Key Innovations

- **Progressive Growing:**
 - Starts with low-resolution images
 - Gradually increases resolution
 - Preserves high-frequency details
- **Adaptive Instance Normalization:**
 - Style vectors control image appearance
 - Enables style mixing experiments
- **No Batch Normalization:**
 - Uses instance normalization instead
 - Better preserves individual image styles

8.2.2 Improved GAN Training

Crucial Insight

Modern GAN Techniques

- **Wasserstein GAN with Gradient Penalty:**
 - Addresses mode collapse
 - Uses Earth-Mover distance
 - Adds gradient penalty to stabilize training
- **Spectral Normalization:**
 - Stabilizes discriminator training
 - Prevents exploding gradients
 - Improves training convergence
- **Self-Attention GANs:**
 - Adds attention mechanisms
 - Better captures long-range dependencies
 - Improves image coherence

9 Conclusion

9.1 Key Takeaways

Key Takeaway

Architectural Choices Matter

- **Diffusion models** excel in:
 - High-fidelity image generation
 - Text-to-image alignment
 - Controllable generation
- **GANs** excel in:
 - Fast generation
 - Photorealism
 - Style transfer

Key Takeaway

Tool Selection Criteria

- **For educational use:** Learning Assistant (GPT-5) with domain specialization
- **For creative tasks:** Qwen Chat with multimodal capabilities
- **For research:** Qwen's open models with RAG capabilities

9.2 Future Directions

Crucial Insight

Emerging Trends

- **Unified Architectures:**
 - Models combining diffusion and GAN approaches
 - Example: DiffusionGAN for controlled generation
- **Multimodal Diffusion:**
 - Extending diffusion to video and audio
 - Example: Stable Video Diffusion
- **Neural Rendering:**
 - Combining diffusion with 3D reconstruction
 - Example: 3D-Aware Diffusion

9.3 Responsible AI Considerations

Remember

Implementation Best Practices

- **Data Privacy:**
 - Consider on-premise deployment for sensitive data
 - Implement differential privacy for training data
- **Model Safety:**
 - Implement robust guardrails
 - Use intent detection systems
 - Regular safety testing
- **Transparency:**
 - Document model limitations
 - Provide clear citations for generated content
 - Implement explainability features